Highlights

Resource-Bounded Non-Parametric Bootstrap Control Chart for High-Dimensional Non-Gaussian Data Streams via RBULT

P. Junsawang, S. Phimoltares, C. Lursinsap

- We propose **RBULT-SPC**, a novel resource-bounded non-parametric control chart framework for high-dimensional non-Gaussian data streams.
- Combines foundational 1D probabilistic chunk excursion proofs with multivariate Bonferroni Family-Wise Error Rate (FWER) tail scaling.
- Achieves strict $O(D)$ constant RAM memory boundedness (3.23 KB), delivering a 180× memory reduction over conventional sliding-window bootstrap (582.87 KB).
- Integrates Algorithm 4 Z-score spike filtering with dynamic 11-candidate distribution MLE fitting to prevent control limit pollution.
- Eliminates batch false-alarm spam (Chunk FAR = 0.00%) on high-dimensional streams (D = 34) while retaining an immediate anomaly detection response delay (ARL_1 = 1.00).
- Validated across a comprehensive two-tier benchmark: 10 simulated 1D distributions, 4 economic datasets, 5 real industrial streams, and 4 TEP stress regimes (Modes 1, 3, 4, 5) up to 1.74M observations.

Resource-Bounded Non-Parametric Bootstrap Control Chart for High-Dimensional Non-Gaussian Data Streams via RBULT★

P. Junsawang[a], S. Phimoltares[a,*] and C. Lursinsap[a]

[a]*Advanced Virtual and Intelligent Computing (AVIC) Research Center, Department of Mathematics and Computer Science, Faculty of Science, Chulalongkorn University, Bangkok, 10330, Thailand*

ARTICLE INFO

Keywords:
Online Bootstrapping
Statistical Process Control
High-Dimensional IoT Data Streams
Non-Gaussian Resampling
Limited Memory Environments

ABSTRACT

Bootstrapping is a powerful non-parametric technique for quantifying uncertainty without restrictive distributional assumptions. However, applying conventional bootstrap methods to real-time industrial data streams poses two fundamental challenges: prohibitive memory overhead ($O(N \cdot D)$ storage growth) and susceptibility to control limit pollution when anomalous samples contaminate historical memory buffers. This paper presents **RBULT-SPC**, a resource-bounded non-parametric control chart framework designed for high-dimensional, non-Gaussian IoT data streams. Grounded in mathematical proofs of chunk extreme probabilities, RBULT-SPC maintains only compact tail summary statistics in $O(D)$ constant space, preventing memory overflow and eliminating buffer pollution. To handle high dimensionality ($D = 34$), the framework combines Z-score outlier filtering, lazy boundary expansion, 11-candidate non-parametric distribution MLE fitting, and Bonferroni Family-Wise Error Rate (FWER) tail adjustment. Extensive two-tier empirical evaluations across 10 simulated 1D distributions, 4 economic datasets, 5 industrial streams, and 4 operating regimes of the Tennessee Eastman Process (TEP) demonstrate that RBULT-SPC achieves optimal coverage (93.71%−99.95%) and controlled false alarm rates (0.05%−6.29%), outperforming classical Shewhart and EWMA charts whose false alarm rates collapse to 18.21% − 39.01% under non-Gaussianity and operational throughput stress. Furthermore, hyperparameter sensitivity analysis demonstrates that dimension-aware thresholding completely eliminates batch false alarms (Chunk FAR = 0.00%) while maintaining an immediate failure detection response delay ($\mathrm{ARL}_1 = 1.00$) at an average latency under 45 ms.

1. Introduction

In modern smart factory environments and Industry 4.0 applications, real-time quality control relies heavily on monitoring continuous sensor telemetry (e.g., pressure, temperature, current, vibration, rotational speed). Statistical Process Control (SPC) charts serve as the foundational mechanism for detecting out-of-control (OOC) states and operational faults.

1.1. Industrial Context & Problem Motivation

Classical SPC techniques, including Shewhart $\bar{X}$-charts, EWMA, and CUSUM, rely on the parametric assumption that underlying telemetry streams follow a Gaussian normal distribution. However,

★This work was supported in part by the Advanced Virtual and Intelligent Computing (AVIC) Research Center, Department of Mathematics and Computer Science, Faculty of Science, Chulalongkorn University.

*Corresponding author

jpremlc@gmail.com (P. Junsawang); suphakant.p@chula.ac.th (S. Phimoltares); chidchanok.L@chula.ac.th (C. Lursinsap)

ORCID(s):

real-world sensor streams in industrial processes frequently exhibit non-Gaussian characteristics, including extreme skewness, heavy tails, and negative kurtosis.

1.2. Critical Bottlenecks of Existing Methods

Applying traditional resampling and bootstrap methods to high-dimensional streaming environments encounters three primary bottlenecks:

1. **Memory Storage Explosion ($O(N \cdot D)$):** Storing historical observations indefinitely leads to Out-Of-Memory (OOM) failures on resource-constrained embedded IoT edge devices.
2. **Control Limit Pollution:** When anomalous fault observations contaminate sliding memory buffers, conventional bootstrap limits expand unnaturally, masking subsequent anomalies.
3. **High-Dimensional False Alarm Spam:** Monitoring D sensor channels simultaneously without Family-Wise Error Rate (FWER) adjustment results in an unacceptable rate of false alarms.

1.3. Summary of Contributions

To address these challenges, this paper presents **RBULT-SPC**. The major contributions of this work include:

- Mathematical formulation and proofs of univariate chunk extreme occurrence probabilities (Theorems 1–2, Corollary 1) and multivariate FWER control (Theorem 3).
- A strict $O(D)$ constant RAM architecture (3.23 KB), achieving a 180× RAM reduction over sliding-window bootstrap (582.87 KB).
- A 4-stage statistical pipeline integrating Z-score spike filtering, lazy boundary expansion, 11-candidate distribution MLE fitting, and Bonferroni tail adjustment.
- Extensive two-tier empirical validation across 10 synthetic 1D distributions, 4 economic datasets, 5 industrial telemetry streams, and 4 TEP operating regimes ($D = 34$, up to 1.74M observations).

2. Related Work

This section reviews relevant literature across three domain areas:

2.1. Classical and Multivariate Statistical Process Control

Review of Shewhart, EWMA, CUSUM, and Hotelling's T^2 control charts, highlighting performance degradation under non-Gaussianity.

2.2. Bootstrap Resampling in Streaming Environments

Review of Efron's percentile bootstrap, blockwise bootstrap, and online streaming bootstrap implementations, noting memory overhead and buffer pollution risks.

2.3. Uncertainty Quantification and Outlier Filtering in IoT Data Streams

Review of streaming range estimation, hyper-rectangle bounding boxes, and Z-score outlier filtering techniques.

3. Theoretical Foundations and Univariate Range Estimation

This section presents the mathematical foundation and 1D range estimation formulation of the proposed Recursively Bootstrapping at Upper-Lower Tails (RBULT) concept, establishing the theoretical ground truth for streaming boundary expansion.

3.1. Concerned Scenario

Sampled data continuously flow in forms of chunks into a process, which may be a learning process or a statistical analyzing process. The size of each chunk is not fixed and does not exceed the size of the working memory. Each sampled chunk contains a set of positive numbers whose minimum and maximum values are unknown. The probability distribution of each integer or continuous value is not predefined and unknown. This scenario exists in various problems in machine learning and data stream analytics. Estimating the population range or variance of incoming data in advance under the *discard-after-learn concept* helps establish adaptive boundaries without retaining past raw observations.

3.2. Studied Problems

Suppose the size of working memory is fixed at m. Let c_i be the i^{th} incoming numeric data chunk such that $|c_i| < m$ and $i \geq 1$. After receiving chunk c_i at time $t = i$, the minimum and maximum values of population at the left and right ends are estimated. Let L denote the estimated left-end value which may be the minimum data value of chunk c_i or the minimum value computed from bootstrapping data in c_i, depending on which one is smaller. Similar to L, let R denote the estimated right-end value which may be the maximum data value of chunk c_i or the maximum value computed from bootstrapping data in c_i, depending on which one is larger. After estimating L and R of c_i, all data except partial summary statistics at both ends are discarded.

1. How to estimate the population variance interval without using the entire data in working memory set D so that data in subsequent incoming chunks fall within the new estimated interval $[L, R]$?
2. Theoretically, how many times on average must the data interval in set D be expanded according to the estimated population variance?

3.3. Constraints

Let $|c_i|$ represent the number of elements in a given chunk c_i. The bin set B is an interval with a left end L and a right end R. The following constraints apply:

1. The probability distribution of data within a chunk c_i is unknown.
2. The size of chunk $|c_i|$ may vary across streaming iterations, with $|c_i| < |D|$.
3. Once all data in chunk c_i have been evaluated, the raw data chunk is completely discarded and cannot be re-inspected.

3.4. Concept of Recursively Bootstrapping at Upper-Lower Tails (RBULT)

RBULT resamples data only in the left-end and right-end tail intervals (b_1 and b_8), differing from classical bootstrap which resamples across the entire data body. The size of left-end and right-end intervals is determined from the standard histogram shape of the best-fitting probability distribution. If the heights of left-end ($|b_1|$) or right-end ($|b_8|$) intervals in the actual incoming data histogram exceed the theoretical heights ($|h_1|$ or $|h_8|$) of the standard histogram, the boundary intervals are extended by bootstrapping data in b_1 and b_8.

The RBULT execution operates in two phases:

1. **Phase 1 (Initial Fitting):** Analyze the probability distribution of the first chunk c_1 and compare the standard histogram with the actual histogram. If heights exceed standard thresholds, bootstrap b_1 and b_8 to extend bounds $[L, R]$.
2. **Phase 2 (Streaming Recurrence):** For each new incoming chunk c_i ($i > 1$), update interval bins, evaluate boundary excursions, and recursively expand $[L, R]$ whenever actual tail heights exceed standard thresholds. Discard non-tail data afterwards.

3.5. Candidate Distribution Selection and CDF Tail Bin Area Integration

To capture diverse physical telemetry behaviors in smart manufacturing (e.g., right-skewed bearing friction, Weibull mechanical wear, Wald diffusion processes, and negative-kurtosis uniform signals), RBULT evaluates a candidate family of continuous distributions $\mathcal{P} = \{\text{Gamma, ExponWeibull, Weibull}_{\text{min / max}}, \text{Wald, ExponPow, Rayleigh, Powerlaw, Lognormal, Chi-Square, Uniform, Normal}\}$. Each candidate represents distinct tail skewness and kurtosis characteristics observed in real-world IoT sensor streams. To construct the standard histogram $H = \{h_1, \ldots, h_8\}$ without relying on standard normal assumptions, the candidate family of continuous distributions is evaluated on initial data chunk c_1. The candidate distribution achieving the optimal parameter fit $\hat{\boldsymbol{\theta}}$ is then selected by minimizing the Akaike Information Criterion (AIC), defined as:

$$\text{AIC} = 2k - 2\ln \mathcal{L}(\hat{\boldsymbol{\theta}} \mid c_1) \tag{1}$$

where k represents the number of estimated parameters (complexity penalty) and $\ln \mathcal{L}(\hat{\boldsymbol{\theta}} \mid c_1)$ denotes the maximized log-likelihood value.

For each candidate density $f(x; \boldsymbol{\theta}) \in \mathcal{P}$, the parameter vector $\boldsymbol{\theta}$ is estimated via Maximum Likelihood Estimation (MLE) by maximizing the log-likelihood function:

$$\hat{\boldsymbol{\theta}} = \arg\max_{\boldsymbol{\theta}} \ln \mathcal{L}(\boldsymbol{\theta} \mid c_1) = \arg\max_{\boldsymbol{\theta}} \sum_{j=1}^{|c_1|} \ln f(x_{1,j}; \boldsymbol{\theta}) \tag{2}$$

The candidate distribution achieving the optimal parameter fit $\hat{\boldsymbol{\theta}}$ is then selected by minimizing the Akaike Information Criterion (AIC), defined as:

$$\text{AIC} = 2k - 2\ln \mathcal{L}(\hat{\boldsymbol{\theta}} \mid c_1) \tag{3}$$

where k represents the number of estimated parameters (complexity penalty) and $\ln \mathcal{L}(\hat{\boldsymbol{\theta}} \mid c_1)$ denotes the maximized log-likelihood value. Minimizing AIC ($\arg\min \text{AIC}$) selects the optimal candidate distribution $f(x; \hat{\boldsymbol{\theta}}) \in \mathcal{P}$ that achieves the best trade-off between fitting accuracy and model simplicity (parsimony). This prevents selecting unnecessarily complex distributions that might overfit transient noise spikes in initial data chunk c_1. Once $\hat{\boldsymbol{\theta}}$ is selected, the expected theoretical element count $|h_k|$ for each bin b_k ($k \in \{1, \ldots, 8\}$) is computed via exact CDF integration:

$$|h_1| = |c_1| \cdot \int_{\mu-4\sigma}^{\mu-3\sigma} f(x; \hat{\boldsymbol{\theta}})\, dx = |c_1| \left[F(\mu - 3\sigma; \hat{\boldsymbol{\theta}}) - F(\mu - 4\sigma; \hat{\boldsymbol{\theta}}) \right] \tag{4}$$

$$|h_8| = |c_1| \cdot \int_{\mu+3\sigma}^{\mu+4\sigma} f(x; \hat{\boldsymbol{\theta}})\, dx = |c_1| \left[F(\mu + 4\sigma; \hat{\boldsymbol{\theta}}) - F(\mu + 3\sigma; \hat{\boldsymbol{\theta}}) \right] \tag{5}$$

If the actual observed data counts $|b_1|$ or $|b_8|$ in chunk c_1 exceed $|h_1|$ or $|h_8|$, the boundary expansion loop (Algorithm 3) is triggered.

3.6. RBULT Algorithms

This subsection details the 4 algorithmic specifications governing 1D RBULT streaming boundary estimation, dual-protocol tracking, and Z-score spike filtering.

Algorithm 1 Initial Chunk Fitting and Tail Summary Setup (c_1)

Require: First data chunk c_1, List $\mathcal{P}$ of candidate probability functions, Tail interval set $B = \{b_1, \dots, b_8\}$, Standard histogram $H = \{h_1, \dots, h_8\}$.
Ensure: Initial left-end bound L_1 and right-end bound R_1.

Load c_1 into working buffer D.
Estimate parameters $\hat{\theta}$ via MLE and select best distribution $f(x; \hat{\theta}) \in \mathcal{P}$ minimizing AIC/KS.
Compute expected tail counts $|h_1|$ and $|h_8|$ via exact CDF integration over $[\mu - 4\sigma, \mu - 3\sigma]$ and $[\mu + 3\sigma, \mu + 4\sigma]$.
Initialize $L_1 = \min(D)$ and $R_1 = \max(D)$.
while $|b_1| > |h_1|$ **or** $|b_8| > |h_8|$ **do**
    Bootstrap tail intervals b_1 and b_8 via Algorithm 3 to refine L_1 and R_1.
    Update bin center $\mu = \frac{L_1+R_1}{2}$ and width $\sigma = \frac{R_1-L_1}{8}$, and recompute B and H.
end while
Store tail summary statistics in b_1, b_8 and discard raw chunk data from D.

Algorithm 2 Streaming Recurrent Boundary Update with Dual Protocol Tracking ($c_m, m > 1$)

Require: Streaming chunk c_m, Existing bounds L_{m-1}, R_{m-1}, Tail set B, Outlier threshold τ.
Ensure: Updated bounds L_m, R_m, Pre-Sequential bounds $[L_m^{\text{pre}}, R_m^{\text{pre}}]$, In-Sample bounds $[L_m^{\text{in}}, R_m^{\text{in}}]$.

while new streaming chunk c_m arrives **do**
    Load c_m into working buffer D.
    Record Pre-Sequential Bounds (Pre-Update Predictive Bounds):
        Set $[L_m^{\text{pre}}, R_m^{\text{pre}}] = [L_{m-1}, R_{m-1}]$.
    Filter transient outlier spikes from D via Algorithm 4 using threshold τ.
    Lazy Boundary Expansion:
    if $\min(D) < L_{m-1}$ **then**
        $L_m = \min(D)$.
    else
        $L_m = L_{m-1}$.
    end if
    if $\max(D) > R_{m-1}$ **then**
        $R_m = \max(D)$.
    else
        $R_m = R_{m-1}$.
    end if
    Reconstruct data histogram B and standard histogram H using $\mu = \frac{L_m+R_m}{2}$ and $\sigma = \frac{R_m-L_m}{8}$.
    while $|b_1| > |h_1|$ **or** $|b_8| > |h_8|$ **do**
        Resample tail bins b_1 and b_8 to expand L_m and R_m via Algorithm 3.
        Recompute μ, σ and update histograms B and H.
    end while
    Record In-Sample Bounds (Post-Update Adapted Bounds):
        Set $[L_m^{\text{in}}, R_m^{\text{in}}] = [L_m, R_m]$.
    Discard raw stream samples D to preserve $O(1)$ constant RAM storage.
end while

Algorithm 3 Recurrent Tail Resampling for Boundaries L and R

Require: Tail data intervals b_1 and b_8, Bootstrap sample size N_{boot}, Iterations M_{iter}.
Ensure: Refined left bound L and right bound R.

— Left Tail Boundary Expansion (L via b_1) —
Initialize mean set $\mathcal{M}_{b_1} = \emptyset$ and std set $\mathcal{S}_{b_1} = \emptyset$.
for $k = 1$ **to** M_{iter} **do**
  Sample N_{boot} points from b_1 with replacement; calculate sample mean $\hat{\mu}_{k,b_1}$ and std $\hat{\sigma}_{k,b_1}$.
  Update $\mathcal{M}_{b_1} = \mathcal{M}_{b_1} \cup \{\hat{\mu}_{k,b_1}\}$ and $\mathcal{S}_{b_1} = \mathcal{S}_{b_1} \cup \{\hat{\sigma}_{k,b_1}\}$.
end for
Compute bootstrap mean $\mu_{b_1}^{(\text{boot})} = \text{mean}(\mathcal{M}_{b_1})$ and dispersion $\sigma_{b_1}^{(\text{boot})} = \text{mean}(\mathcal{S}_{b_1})$.
if $\text{mean}(b_1) < \mu_{b_1}^{(\text{boot})}$ **then**
  $L = \min(b_1) - |\text{std}(b_1) - \sigma_{b_1}^{(\text{boot})}|$.
else
  $L = \min(b_1) - |\text{mean}(b_1) - \mu_{b_1}^{(\text{boot})}|$.
end if
— Right Tail Boundary Expansion (R via b_8) —
Initialize mean set $\mathcal{M}_{b_8} = \emptyset$ and std set $\mathcal{S}_{b_8} = \emptyset$.
for $k = 1$ **to** M_{iter} **do**
  Sample N_{boot} points from b_8 with replacement; calculate sample mean $\hat{\mu}_{k,b_8}$ and std $\hat{\sigma}_{k,b_8}$.
  Update $\mathcal{M}_{b_8} = \mathcal{M}_{b_8} \cup \{\hat{\mu}_{k,b_8}\}$ and $\mathcal{S}_{b_8} = \mathcal{S}_{b_8} \cup \{\hat{\sigma}_{k,b_8}\}$.
end for
Compute bootstrap mean $\mu_{b_8}^{(\text{boot})} = \text{mean}(\mathcal{M}_{b_8})$ and dispersion $\sigma_{b_8}^{(\text{boot})} = \text{mean}(\mathcal{S}_{b_8})$.
if $\text{mean}(b_8) > \mu_{b_8}^{(\text{boot})}$ **then**
  $R = \max(b_8) + |\text{std}(b_8) - \sigma_{b_8}^{(\text{boot})}|$.
else
  $R = \max(b_8) + |\text{mean}(b_8) - \mu_{b_8}^{(\text{boot})}|$.
end if

Algorithm 4 Z-Score Outlier Spike Filtering for Contaminated Streams

Require: Data buffer D, Spike sensitivity threshold τ (default $\tau = 3.0$).
Ensure: Filtered data buffer D free of transient noise spikes.

Compute sample mean $\mu_D = \frac{1}{|D|}\sum_{x \in D} x$ and sample std $\sigma_D = \sqrt{\frac{1}{|D|-1}\sum_{x \in D}(x - \mu_D)^2}$.
for each sample $x \in D$ **do**
  Calculate standard score $z(x) = \frac{|x - \mu_D|}{\sigma_D}$.
  if $z(x) > \tau$ **then**
    Exclude outlier spike x from buffer D.
  end if
end for
Return cleaned working buffer D.

3.7. Theoretical Analysis

The theoretical minimum and maximum numbers of boundary expansion can be analyzed in terms of probability combinatorics. Let population size be N and chunk size be l without duplicate observations. Let $p^{(\min)} = \min(\mathcal{P})$ and $p^{(\max)} = \max(\mathcal{P})$.

Theorem 1 (First Chunk Extreme Probability)**.** *Assume a chunk of l observations is sampled without replacement from population $\mathcal{P}$ of size N. The probability of containing both $p^{(\min)}$ and $p^{(\max)}$ in the first chunk is:*

$$P(E_1) = \frac{l(l-1)}{N(N-1)} \tag{6}$$

Proof. When $p^{(\min)}$ and $p^{(\max)}$ are present in the first chunk, the remaining $l-2$ observations must be chosen from the remaining $N-2$ population elements. The number of ways to pick $l-2$ elements from $N-2$ is $\binom{N-2}{l-2}$. Combining with $(N-l)!$ permutations of remaining population items yields $c = \binom{N-2}{l-2}(N-l)! = \frac{(N-2)!}{(l-2)!}$. Incorporating $l!$ chunk permutations yields $l(l-1)(N-2)!$ total favorable permutations out of $N!$ possible population arrangements. Thus:

$$P(E_1) = \frac{l(l-1)(N-2)!}{N!} = \frac{l(l-1)}{N(N-1)} \tag{7}$$

□

Theorem 2 (Streaming Chunk Distribution)**.** *Assume a chunk of l observations is sampled without replacement from population $\mathcal{P}$ of size N. The probability of having $p^{(\min)}$ and $p^{(\max)}$ in the i^{th} incoming chunk ($i \geq 1$) is uniform and equal to:*

$$P(E_i) = \frac{l(l-1)}{N(N-1)} \tag{8}$$

Proof. Let S denote the sample space of partitioning N observations into L chunks of size l ($N = l \cdot L$). The total number of sample outcomes is $|S| = \frac{N!}{(l!)^L}$. Let $E_i \subset S$ denote events where both extreme bounds $p^{(\min)}$ and $p^{(\max)}$ reside in chunk i. By product rule combinatorics:

$$|E_i| = \binom{N-2}{l-2}\frac{(N-l)!}{(l!)^{L-1}} \tag{9}$$

Evaluating the probability ratio yields:

$$P(E_i) = \frac{|E_i|}{|S|} = \frac{\binom{N-2}{l-2}\frac{(N-l)!}{(l!)^{L-1}}}{\frac{N!}{(l!)^L}} = \frac{l(l-1)}{N(N-1)} \tag{10}$$

which proves uniform occurrence probability across all streaming chunk positions $i \geq 1$. □

Corollary 1 (Expected Chunk Index for Extreme Bounds)**.** *Assuming joint occurrence of minimum and maximum observations is restricted to a single chunk, the expected chunk index $\mathbb{E}[X]$ containing both extreme bounds across L total chunks is:*

$$\mathbb{E}[X] = \frac{1}{2}\left(\frac{l(l-1)L(L+1)}{N(N-1)}\right) \tag{11}$$

Proof. Let $X \in \{1, 2, \dots, L\}$ be a discrete random variable representing the chunk index containing $p^{(\min)}$ and $p^{(\max)}$. Summing over uniform probability $P(E_i) = \frac{l(l-1)}{N(N-1)}$ yields:

$$\begin{aligned}\mathbb{E}[X] &= \sum_{i=1}^{L} i \cdot P(X=i) = \left(\sum_{i=1}^{L} i\right)\frac{l(l-1)}{N(N-1)} \\ &= \left(\frac{L(L+1)}{2}\right)\left(\frac{l(l-1)}{N(N-1)}\right) = \frac{1}{2}\left(\frac{l(l-1)L(L+1)}{N(N-1)}\right)\end{aligned} \tag{12}$$

□

4. Proposed High-Dimensional RBULT-SPC Framework

This section presents the extension of the 1D RBULT concept to multivariate high-dimensional Statistical Process Control (RBULT-SPC).

4.1. System Architecture and 4-Stage Streaming Pipeline

Details of the 4-stage streaming process: Stationary Preprocessing, Z-score Spike Filter, Dimensional Bound Estimation, and Bounding Hyper-Rectangle $\mathcal{B}_m = \prod_{d=1}^{D}[L_{m,d}, R_{m,d}]$ evaluation.

4.2. Bonferroni / idák Family-Wise Error Rate (FWER) Scaling

To maintain a target system false alarm rate $\alpha_{\text{sys}} = 0.05$ across D monitored sensor channels, per-dimension tail quantiles are scaled using Bonferroni correction:

$$\alpha_{\text{dim}} = \frac{\alpha_{\text{sys}}}{D} \tag{13}$$

yielding dimension-aware control limit formulations:

$$\text{LCL}_d = L_d = Q_{\text{tail}}\left(\frac{\alpha_{\text{sys}}}{2D}\right), \quad \text{UCL}_d = R_d = Q_{\text{tail}}\left(1 - \frac{\alpha_{\text{sys}}}{2D}\right) \tag{14}$$

4.3. Theorem 3 (Family-Wise Error Rate Upper Bound)

Theorem 3 (Multivariate FWER Upper Bound)**.** *For a D-dimensional streaming process monitored by bounding hyper-rectangle $\mathcal{B}_m = \prod_{d=1}^{D}[L_{m,d}, R_{m,d}]$ with per-channel Bonferroni tail scaling $\alpha_{dim} = \frac{\alpha_{sys}}{D}$, the overall system false alarm probability under in-control conditions is upper-bounded by α_{sys}:*

$$P\left(\exists d \in \{1, \dots, D\} : x_d \notin [L_d, R_d] \mid \textit{In-Control}\right) \leq \alpha_{sys} \tag{15}$$

Proof. Let E_d denote the event that the d^{th} sensor dimension exceeds its marginal control boundaries under in-control conditions:

$$E_d = \{x_d \notin [L_d, R_d] \mid \text{In-Control}\}, \quad \text{for } d = 1, \dots, D \tag{16}$$

By definition of the per-channel control limits $[L_d, R_d]$ using Bonferroni-scaled quantiles $\alpha_{\text{dim}} = \frac{\alpha_{\text{sys}}}{D}$, the marginal violation probability per channel is:

$$\begin{aligned} P(E_d) &= P(x_d < L_d) + P(x_d > R_d) \\ &= \frac{\alpha_{\text{sys}}}{2D} + \frac{\alpha_{\text{sys}}}{2D} = \frac{\alpha_{\text{sys}}}{D} \end{aligned} \tag{17}$$

A system-wide false alarm occurs if at least one sensor dimension experiences a boundary violation, corresponding to the union event $\bigcup_{d=1}^{D} E_d$. By Boole's inequality (the union bound), the joint probability of the union of arbitrary events is strictly upper-bounded by the sum of their marginal probabilities, regardless of the underlying cross-channel correlation structure $\mathbf{\Sigma}$:

$$\begin{aligned} P\left(\exists d \in \{1, \dots, D\} : x_d \notin [L_d, R_d] \mid \text{In-Control}\right) &= P\left(\bigcup_{d=1}^{D} E_d\right) \\ &\leq \sum_{d=1}^{D} P(E_d) \end{aligned}$$

$$= \sum_{d=1}^{D} \frac{\alpha_{\text{sys}}}{D} = D \cdot \frac{\alpha_{\text{sys}}}{D} = \alpha_{\text{sys}} \tag{18}$$

Furthermore, if the D channels are mutually independent, the exact Family-Wise Error Rate simplifies via idák's identity:

$$\text{FWER}_{\text{indep}} = 1 - \prod_{d=1}^{D} \left(1 - P(E_d)\right) = 1 - \left(1 - \frac{\alpha_{\text{sys}}}{D}\right)^{D} < \alpha_{\text{sys}} \tag{19}$$

Since $\text{FWER}_{\text{indep}} < \sum_{d=1}^{D} P(E_d) = \alpha_{\text{sys}}$, Boole's inequality guarantees that α_{sys} remains a strict upper bound under arbitrary sensor cross-correlations, completing the proof. □

4.4. Asymptotic Complexity Analysis

- **Space Complexity:** Strictly $O(D)$ bounded RAM (3.23 KB for $D = 34$), independent of stream length N.
- **Time Complexity:** Amortized $O(1)$ per sample (< 45 ms per streaming batch).

5. Two-Tier Empirical Benchmark Experiments

This section presents empirical evaluation results across a two-tier benchmark suite.

5.1. Experimental Setup & Evaluation Metrics

Formulation of Coverage Rate (%), Sample-level FAR (%), Chunk-level FAR (%), ARL_0, ARL_1 (Detection Delay), Peak Memory Footprint (KB), and Avg Streaming Latency (ms).

5.2. Tier 1: Univariate Range Approximation Benchmarks

Performance on 10 simulated 1D distributions ($F, Wald, Normal, \chi^2, Weibull$) and 4 real-world economic datasets.

5.3. Tier 2: Real-World Industrial Telemetry Benchmarks

Benchmark results across AI4I 2020 Predictive Maintenance ($N = 10k, D = 5$), MetroPT-3 Air Compressor ($N = 1.51M, D = 7$), Large Industrial Pump ($N = 20k, D = 5$), and Water Pump Sensor ($N = 220k, D = 10$).

5.4. Tier 2 Stress Benchmark: Tennessee Eastman Process (TEP Multi-Mode)

Comparative performance across TEP Modes 1 (Nominal), 3 (Feed Skewness 90/10), 4 (Max Production Throughput), and 5 (Combined Extreme Stress) for $D = 34$ sensor channels.

5.5. Hyperparameter Sensitivity Analysis

Evaluation of chunk alarm threshold $C_{\text{thresh}} \in \{5, 10, 15\}$ on TEP Mode 1 stream.

6. Discussion and Data Taxonomy Insights

Comparative discussion of non-Gaussian tail fitting, memory explosion prevention, EWMA lag suppression, and edge IoT deployment guidelines.

7. Conclusion

Summary of key scientific findings and roadmap for Paper 2 (concept drift adaptation).